Lecture 17: Joins

STA courses next semester

- STA 214: Applied GLMs (requires MTH 111)
- STA 310: Probability (requires MTH 112)
- STA 311: Inference (requires 310)
- STA 312: Linear models (requires 310 + MTH 121 or 205)
- STA 352: Networks (requires MTH 117, 121, or 205)
- STA 362: Multivariate (requires MTH 121 or 205)
- STA 363: Stat learning (requires MTH 121 or 205)
- STA 365: Bayesian (requires 310)
- STA 368: Time series (requires 310)

Data stored in multiple tables

The `nycflights13` package contains information on flights from NYC airports in 2013. The data is stored across several data frames:

- `airlines`: information on each airline
- `airports`: information on each airport
- `flights`: information on each flight
- `planes`: information on each plane
- `weather`: hourly weather data

Question: What is the advantage of storing this data in multiple tables, instead of one BIG table?

Data stored in multiple tables

- Databases often contain different tables to store different information
- For example, a healthcare database could contain the following tables:
  - `patients`
  - `doctors`
  - `offices`
  - `insurance`

Joining tables

```
1 patients
```

```
  age insurance  provider
1  23     Aetna Dr. Zhang
2  47      BCBS Dr. Foyle
3  38  Medicaid Dr. Zhang
```

```
1 doctors
```

```
    provider      location
1  Dr. Foyle    Greensboro
2 Dr. Renard Winston-Salem
3  Dr. Zhang Winston-Salem
```

I want to add location information to the `patient` table.

What should the resulting table look like?

Left join

```
1 patients
```

```
  age insurance  provider
1  23     Aetna Dr. Zhang
2  47      BCBS Dr. Foyle
3  38  Medicaid Dr. Zhang
```

```
1 doctors
```

```
     provider      location
1  Dr. Foyle    Greensboro
2 Dr. Renard Winston-Salem
3  Dr. Zhang Winston-Salem
```

```
1 patients |>
2   left_join(doctors, join_by(provider))
```

```
  age insurance  provider      location
1  23     Aetna Dr. Zhang Winston-Salem
2  47      BCBS Dr. Foyle    Greensboro
3  38  Medicaid Dr. Zhang Winston-Salem
```

Left join

```
1 patients |>
2   left_join(doctors, join_by(provider))
```

```
  age insurance  provider      location
1  23     Aetna Dr. Zhang Winston-Salem
2  47      BCBS Dr. Foyle    Greensboro
3  38  Medicaid Dr. Zhang Winston-Salem
```

- Left joins are useful for adding additional information to a table
- Left joins (generally) keep the same rows as the initial dataframe (`patients`), and add more columns
- `join_by` specifies how to link the tables

Left joins in Python

```
1 import pandas as pd
2
3 pd.merge(patients, doctors, how = 'left',
4           left_on = 'provider', right_on = 'provider')
```

```
   age insurance   provider       location
0  23.0     Aetna  Dr. Zhang  Winston-Salem
1  47.0      BCBS  Dr. Foyle     Greensboro
2  38.0  Medicaid  Dr. Zhang  Winston-Salem
```

Joining tables

Flights information:

```
# A tibble: 3 × 5
  time_hour           origin dest  tailnum carrier
  <dttm>              <chr>  <chr> <chr>   <chr>
1 2013-01-01 05:00:00 EWR    IAH   N14228  UA
2 2013-01-01 05:00:00 LGA    IAH   N24211  UA
3 2013-01-01 05:00:00 JFK    MIA   N619AA  AA
```

Weather information

```
# A tibble: 3 × 4
  origin time_hour            temp wind_speed
  <chr>  <dttm>               <dbl>      <dbl>
1 EWR    2013-01-01 01:00:00  39.0       10.4
2 EWR    2013-01-01 02:00:00  39.0        8.06
3 EWR    2013-01-01 03:00:00  39.0       11.5
```

Question: What if I want to get information about the weather for each flight?

Left joins

```
1 flights |>
2   left_join(weather, join_by(origin, time_hour))
```

```
# A tibble: 6 × 7
  time_hour           origin dest  tailnum carrier  temp wind_speed
  <dttm>              <chr>  <chr> <chr>   <chr>   <dbl>      <dbl>
1 2013-01-01 05:00:00 EWR    IAH   N14228  UA       39.0       12.7
2 2013-01-01 05:00:00 LGA    IAH   N24211  UA       39.9       15.0
3 2013-01-01 05:00:00 JFK    MIA   N619AA  AA       39.0       15.0
4 2013-01-01 05:00:00 JFK    BQN   N804JB  B6       39.0       15.0
5 2013-01-01 06:00:00 LGA    ATL   N668DN  DL       39.9       16.1
6 2013-01-01 05:00:00 EWR    ORD   N39463  UA       39.0       12.7
```

Joining with different names

Suppose our tables looked like this:

```
1 patients
```

```
  age insurance  provider
1  23     Aetna Dr. Zhang
2  47      BCBS Dr. Foyle
3  38  Medicaid Dr. Zhang
```

```
1 doctors
```

```
         name      location
1  Dr. Foyle    Greensboro
2 Dr. Renard Winston-Salem
3  Dr. Zhang Winston-Salem
```

How would we specify the columns to link the tables?

Joining with different names

Suppose our tables looked like this:

```
1 patients
```

```
  age insurance  provider
1  23     Aetna Dr. Zhang
2  47      BCBS Dr. Foyle
3  38  Medicaid Dr. Zhang
```

```
1 doctors
```

```
         name      location
1  Dr. Foyle    Greensboro
2 Dr. Renard Winston-Salem
3  Dr. Zhang Winston-Salem
```

```
1 patients |>
2   left_join(doctors, join_by(provider == name))
```

```
  age insurance  provider      location
1  23     Aetna Dr. Zhang Winston-Salem
2  47      BCBS Dr. Foyle    Greensboro
3  38  Medicaid Dr. Zhang Winston-Salem
```

In Python

```
1 pd.merge(patients, doctors, how = 'left',
2           left_on = 'provider', right_on = 'name')
```

	age	insurance	provider	name	location
0	23.0	Aetna	Dr. Zhang	Dr. Zhang	Winston-Salem
1	47.0	BCBS	Dr. Foyle	Dr. Foyle	Greensboro
2	38.0	Medicaid	Dr. Zhang	Dr. Zhang	Winston-Salem

Another join

Patients in the system:

```
1 patients
```

	age	insurance	provider
1	23	Aetna	Dr. Zhang
2	47	BCBS	Dr. Foyle
3	38	Medicaid	Dr. Zhang
4	54	Humana	Dr. Renard

Accepted insurance:

```
1 insurance
```

	company	phone
1	Anthem	800-676-2583
2	BCBS	877-258-3334
3	Kaiser	800-810-4766
4	Medicaid	877-201-3750

Suppose I want insurance information only for the patients who have an accepted insurance. What should the final table look like?

Inner join

Patients in the system:

```
1 patients
```

```
  age insurance   provider
1  23     Aetna  Dr. Zhang
2  47      BCBS  Dr. Foyle
3  38  Medicaid  Dr. Zhang
4  54    Humana Dr. Renard
```

Accepted insurance:

```
1 insurance
```

```
   company         phone
1   Anthem 800-676-2583
2     BCBS 877-258-3334
3   Kaiser 800-810-4766
4 Medicaid 877-201-3750
```

```
1 patients |>
2   inner_join(insurance, join_by(insurance == company))
```

```
  age insurance  provider         phone
1  47      BCBS Dr. Foyle 877-258-3334
2  38  Medicaid Dr. Zhang 877-201-3750
```

In Python

```
pd.merge(patients, insurance, how='inner',
         left_on = 'insurance', right_on = 'company')
```

```
    age insurance   provider   company         phone
0  47.0      BCBS  Dr. Foyle      BCBS  877-258-3334
1  38.0  Medicaid  Dr. Zhang  Medicaid  877-201-3750
```

Class activity

https://sta279-f23.github.io/class_activities/ca_lecture_17.html